Databases

>Relational Database Service

RDS offers a managed and highly scalable database environment for most popular relational database engines (including MySQL, MariaDB, and Oracle).

The Amazon Relational Database Service (RDS) is Amazon's managed relational database service. RDS lets you provision a number of popular relational database management systems (RDBMSs) including Microsoft SQL Server, Oracle, MySQL, and PostgreSQL

When you create an RDS database instance, Amazon sets up one or more compute instances and takes care of installing and configuring the RDBMS of your choice.

Automated patching, backups, redundancy, failover and disaster recovery. RDS will automatically fail over to the standby during a failure so database operations can resume quickly without admin intervention.

Joining multiple tables to get the bigger picture. RDS is generally used for online transaction processing (OLTP) workloads

Able to launch read replicas across regions, which is faster for faster querying

>RDS read replicas have their own endpoints

RDS can also perform manual or automatic EBS snapshots that you can easily restore to new RDS instances. RDS can also handle the hard work of installing patches and upgrades during scheduled maintenance windows.

Up and running in minutes, mult-AZ, automated failover capability, automated backups,

A manual install in your own data center could take 8 days or longer. But with the great thing about RDS you can have it up and running in 5-10 mins.

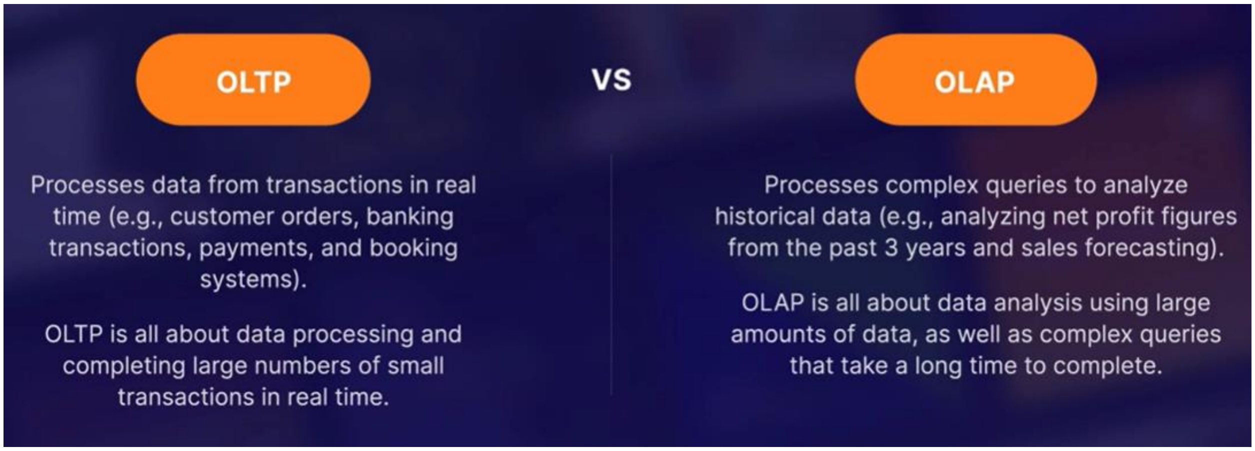

OLTP VS OLAP

OLTP

Processes data from transactions in real time (e.g., customer orders, banking transactions, payments, and booking systems).

OLTP is all about data processing and completing large numbers of small transactions in real time.

OLAP

Processes complex queries to analyze historical data (e.g., analyzing net profit figures from the past 3 years and sales forecasting).

OLAP is all about data analysis using large amounts of data, as well as complex queries that take a long time to complete.

*RDS is not suitable for analyzing large amounts of data. Use a data warehouse like redshift which is optimized for OLAP

Anything about OLAP & OLTP

- Think that online transactions are through RDS (OLTP)
- Product that is called RedShift (OLAP)

Amazon RDS supports the following six database engines:

- MySQL
- MariaDB
- Oracle
- PostgreSQL
- Microsoft SQL Server
- Amazon Aurora

- Aurora
    - With the exception of Amazon Aurora, these database engines are either open source or commercially available products found in many data center environments. Amazon Aurora is a proprietary database designed for RDS, but it's compatible with existing MySQL (5 read), Aurora Replicas (15 read) and PostgreSQL databases (5 read)
    - If you use the Amazon Aurora database engine—Amazon's proprietary database engine designed for and available exclusively with RDS—you can take advantage of additional benefits when using multi-AZ. When you use Aurora, your RDS instances are part of an Aurora cluster.
    - 5x faster than normal MySQL and 3x faster than normal prostgreSQL
    - Managed by RDS
    - Continuous backup to Amazon S3
    - 2 copies of your data are contained in each AZ, with a min of 3 AZ. 6 copies of your data
    - Compute resources can scale up to 96 vCPUs and 768 GB of memory
    - Point-in-time recovery
    - Self healing
    - Aurora stores copies of the data across multiple Availability Zones in a single AWS Region.
    - Aurora Serverless is a supported feature of Aurora.
    - Automatic Backups is a supported feature of Aurora

Aurora Serverless

An Aurora Serverless DB cluster automatically starts up, shuts down, and scales capacity up or down based on your applications needs. Might be some questions of spiky workloads, you want to look at this

Use Cases

- Provides relatively simple, cost-effective option for infrequent, intermittent, or unpredictiable workloads.

- Dynamo DB
    - DynamoDB is Amazon's managed nonrelational database service. It's designed for highly transactional applications that need to read from or write to a database tens of thousands of times a second.
    - analogous to a row or record in a relational database. DynamoDB stores items in tables.
    - Other than the primary key, an item doesn't have to have any particular attributes.

        - This flexibility makes DynamoDB the database of choice for applications that need to store a wide variety of data without having to know the nature of that data in advance.
    - Non-Relational databases such as DynamoDB are designed to scale horizontally by spreading your data across more partitions, allowing for thousands of reads and writes per second
    - DynamoDB stores data as items in tables.
    - Each item must have primary key whose values are unique within the table. This is how DynamoDB uniquely identifies an item.
    - Great for new product launches
    - Pay-per-request pricing, but you pay more per request than with provisioned capacity
    - DynamoDB uses the primary key to distribute items across different partitions. The number of partitions allocated to a table depends on the number of WCU and RCU you configure.
    - DynamoDB is a NoSQL database
    - The minimum monthly availability for DynamoDB is 99.99 percent in a single Region. It's not 99.95 percent, 99.9 percent, or 99.0 percent.
- DynamoDB Accelerator (DAX)
    - Fully managed, highly available, in-memory cache
    - 10x performance improvement
    - No need for developers to manage caching logic (This is done by DynamoDB)
    - Caching with DAX
        - The application speaks directly to DAX to see if the information is cached inside of DAX > If not in cached in DAX it will basically interrogate DynamoDB and bring it up to DAX and back to the application

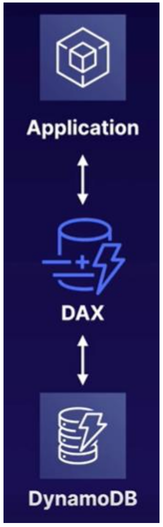

    - Security in DynoDB
        - Encryption at rest using KMS
        - CloudWatch and CloudTraill
        - Site-to-site VPN

- Document database
    - MongoDB compatible
- ElastiCache

    - Fully managed in-memory datastore compatible with redis or memcached
        - Data can be lost because it ‘s stored in-memory
    - High performance and low latency
- Neptune
    - Graph database service
    - Highly connected datasets like social media sites

Down below are what are the best option for each use case.

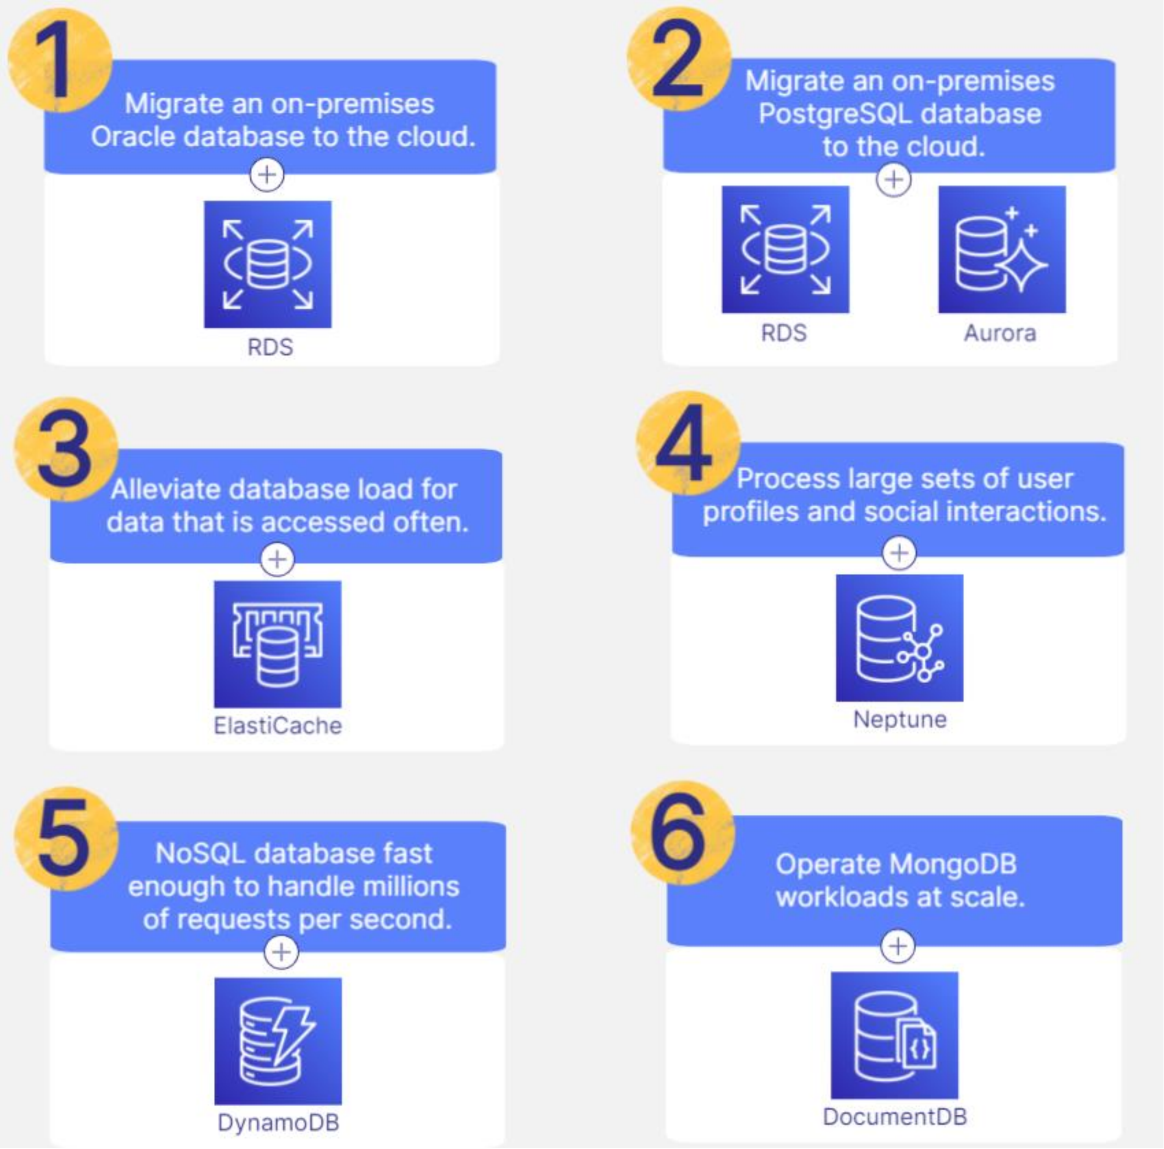

1 Migrate an on-premises Oracle database to the cloud.

2 Migrate an on-premises PostgreSQL database to the cloud.

3 Alleviate database load for data that is accessed often.

4 Process large sets of user profiles and social interactions.

5 NoSQL database fast enough to handle millions of requests per second.

6 Operate MongoDB workloads at scale.

Mult-AZ is for disaster recovery, not for improving performance, so you cannot connect to the standby when the primary database is active.

>Increasing Read Performance with Read Replicas

Read replicas is a read-only copy of your primary database. Which great for read-heavy workloads and takes the load off your primary database. Each replica has its own unique DNS endpoint. This can also be promoted to be their own databases.

*There may be scenarios where they try and trick you in between read replica and Multi-AZ

Multi-AZ is for disaster recovery only, and in the event of a failure, RDS will automatically fail over to the standby instance.

Read replicas are perfect for read-heavy workloads

Read-Replica is for boosting performance, scaling

- Requires automatic backups
- Multiple read replicas are supported. Up to 5 read replicas to each DB instance

> When Do We Use DynamoDB Transactions?

The ACID Diagram

Name	Includes/Function
Atomic	All changes to the data must be performed successfully or not at all.
Consistent	Data must be in a consistent state before and after the transaction
Isolated	No other process can change the data while the transaction is running
Durable	The changes made by a transaction must be persist

You can use ACID with DynamoDB, but you need to use DynamoDB transactions which provides developers atomicity, consistency, isolation and durability.

When using DynamoDB transactions

- Multiple “all-or-nothing” operations
- Financial transactions
- Fullfilling orders
- 3 options for reads: eventual, strong, and transactional consistency
- 2 options for writes: standard and transactional
- Up to 25 items of 4 MB of data

>Savings your data with DynamoDB backups

This on-demand backup and restore can do full backups at any time. Which has ZERO impact on table performance or availability. The consistency is within seconds and retained until deleted. Operates within the same region as the source table

- Point-in-Time recovery (PITR)
    - Protects against accidental writes or deletes
    - Restore to any point in the last 35 days
    - Incremental backups
    - This is not enabled by default
    - Latest restorable: 5 mins in the past

>Taking your data global with DynamoDB Streams and Global Tables

- Streams
  - Time-ordered sequence of item-level changes in a table
  - Stored for 24 hours
  - So every time you make a change to a Dynamo DB table, that data is going to be stored sequentially in a stream record which is broken up into shards
  - Inserts, updates, and deletes
    - This gives us that time-ordered sequence
  - Combine with Lambda functions for functionality like stored procedures
- Global Tables
  - A way to replicate DynamoDB tables from one region to another
  - Based on DynamoDB streams, need this turned on
  - Globally distributed applications
  - Multi-region redundancy for disaster recovery or high availability
  - Replication latency under 1 second

Exam Tips

1. Relational Database Service
   a. 6 RDS Datatypes
   b. RDS is for OLTP Workloads
   c. Not suitable for OLAP workloads
2. Aurora
   a. At least 6 copies of your data
   b. Share snapshots with other AWS accounts
   c. 3 types of replicas: MySQL (5 read), Aurora Replicas (15 read) and PostgreSQL databases (5 read)
   d. Automated backups turned on by default
   e. Aurora serverless if you want a simple, cost-effective option for infrequent, intermittent or unpredictable workloads
3. DynamoDB
   a. Stored on SSD storage
   b. Spread across 3 data centers
   c. Eventually consistent reads (default)
   d. Strongly consistent reads
   e. This on-demand backup and restore can do full backups at any time.

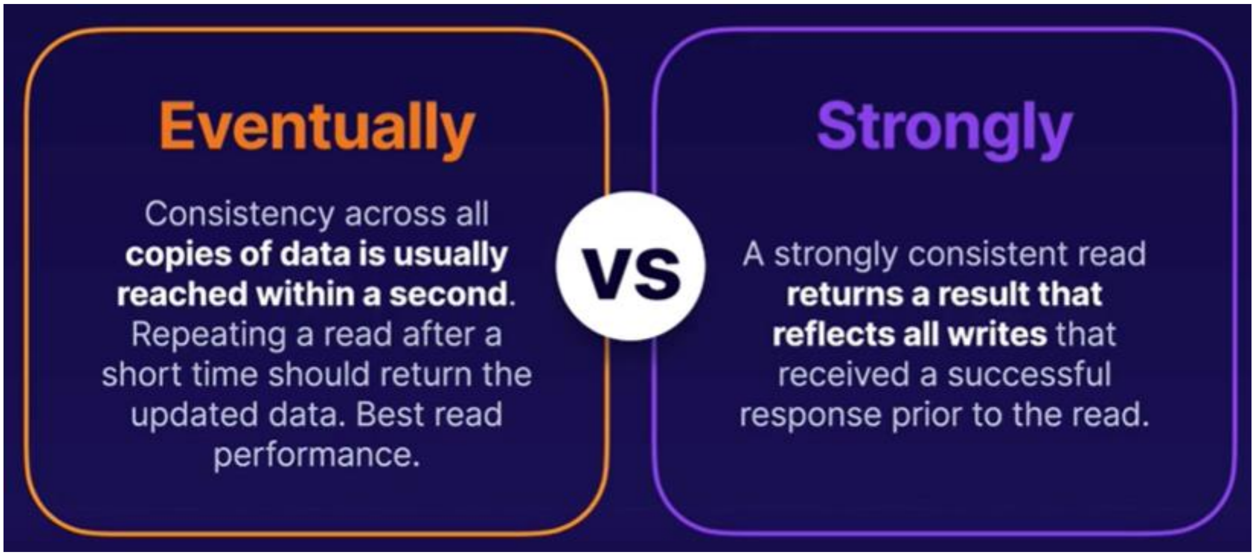

Eventually

Consistency across all **copies of data is usually reached within a second.** Repeating a read after a short time should return the updated data. Best read performance.

VS

Strongly

A strongly consistent read **returns a result that reflects all writes** that received a successful response prior to the read.

4. DynamoDB Transactions
   a. Any scenario questions that mentions ACID requirements, think of this
   b. Transactions provide developers, atomicity, consistency, isolation and durability across 1 or more tables within a single AWS account and region
   c. All-or-nothing transactions

https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadConsistency.html